

## Lab 1 Report

Group Members:

130050007 : Vaibhav Bhosale  
130050046 : Dibyendu Mondal

### Question 1:

Number of cpu cores are : 4  
Total Memory : 8142244 kB  
Memory Free : 1328044 kB  
Fraction Free : 0.16310540435  
Number of context switches : 368384286  
Number of processes forked : 1648009

### Question 2:

For cpu1:

Output of top:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
25872	dibyendu	20	0	4188	356	276	R	100.0	0.0	2:55.49	cpu1

Output of iostat -d -x:

Device:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
sda	0.40	8.40	2.20	22.80	12.80	156.80	13.57	0.19	7.62	14.18	6.98	5.79	14.48

Output of sudo nethogs:

PID	USER	PROGRAM	DEV	SENT	RECEIVED
1809	vpb	/opt/google/chrome/chrome	wlan0	0.594	0.443 KB/sec

The bottleneck resource is CPU.

Since other resources are not being used and CPU is used 100%, it's the bottleneck.

For cpu1printf:

Output of top for the program's process:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
26013	dibyendu	20	0	4192	356	276	S	25.1	0.0	0:01.49	cpu1printf

Output of top for shell process:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
8848	vpb	20	0	623456	29444	20420	R	98.2	0.8	3:02.14	gnome-term+

Output of iostat -d -x:

Device:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
sda	0.00	2.40	0.00	0.40	0.00	11.20	56.00	0.01	28.00	0.00	28.00	28.00	1.12

Output of sudo nethogs:

PID	USER	PROGRAM	DEV	SENT	RECEIVED
1809	vpb	/opt/google/chrome/chrome	wlan0	1.635	1.024 KB/sec

The bottleneck in this case is CPU. But this resource isn't maxed out for this process, but for the shell process which is called in every iteration of the program for printing on the terminal. Thus this program is slowed down because of the shell process whose bottleneck resource is CPU, hence, the bottleneck resource for this program is CPU.

For cpu2:

Output of top:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
26028	dibyendu	20	0	4188	356	276	R	100.0	0.0	0:10.00	cpu2

Output of iostat -d -x:

Device:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
sda	1.20	2.00	0.20	4.20	5.60	26.40	14.55	0.01	2.36	0.00	2.48	2.36	1.04

Output of sudo nethogs:

PID	USER	PROGRAM	DEV	SENT	RECEIVED
1809	vpb	/opt/google/chrome/chrome	wlan0	1.635	1.024 KB/sec

The bottleneck resource is CPU  
Since other resources are not being used and CPU is used 100%, it's the bottleneck.

For disk:

Output of top:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
26093	dibyendu	20	0	4328	528	436	D	1.0	0.0	0:00.07	disk

Output of iostat -d -x:

Device:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
sda	45.60	201.60	522.60	12.60	54481.60	858.40	206.80	2.22	4.16	4.10	6.60	1.86	99.60

Output of sudo nethogs:

PID	USER	PROGRAM	DEV	SENT	RECEIVED
1809	vpb	/opt/google/chrome/chrome	wlan0	1.635	1.024 KB/sec

The bottleneck resource is disk  
%util is 99.6. Since disk I/Os are used, it's the bottleneck.

For disk1:

Output of top:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
26143	dibyendu	20	0	4328	524	432	R	15.9	0.0	0:01.83	disk1

Output of iostat -d -x:

Device:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
sda	53.80	3.80	72.80	0.80	1437.60	18.40	39.57	0.41	5.52	5.20	35.00	2.98	21.92

Output of nethogs:

PID	USER	PROGRAM	DEV	SENT	RECEIVED
1809	vpb	/opt/google/chrome/chrome	wlan0	1.635	1.024 KB/sec

The bottleneck resource is disk  
Since disk I/Os are used, it's the bottleneck.

### Question 3:

For cpu1:

Time in User Mode: 2866  
Time in Kernel Mode: 6

Reason: Time taken in user mode is greater than time taken in kernel mode because the instructions in the program do not need to make any system calls.

For cpu1print:

Time in User Mode: 26  
Time in Kernel Mode: 146

Reason: Time taken in user mode is less than time taken in kernel mode because printf instruction is used which makes system calls resulting in majority time spent to execute the printf function.

For cpu2:

Time in User Mode: 3605

Time in Kernel Mode: 6

Reason: Time taken in user mode is greater than time taken in kernel mode because all instructions except "gettimeofday" do not need to make any system calls. Also, gettimeofday has to just retrieve the system time and hence it is done quite fast.

#### **Question 4:**

For cpu1:

voluntary\_ctxt\_switches: 2

nonvoluntary\_ctxt\_switches: 3443

Reason: The voluntary context switches are less because there are no other events involved due to which the process would very rarely voluntarily give up the CPU. The large number of non-voluntary context switches are due to the infinite loop.

For disk:

voluntary\_ctxt\_switches: 4338

nonvoluntary\_ctxt\_switches: 4841

Reason: The voluntary context switches are more compared to cpu1 because the process voluntarily give up the process to the CPU for disk operations (reading from files). The large number of non-voluntary context switches are due to the infinite loop.